

Final Manifest

CSYE 7220 - Development and Operations

Deep Gamit - 001057437
Priyam Suthar - 001054934

Deploying the Rockstar App to AKS cluster with Prometheus for Monitoring

Step 1: We deploy the resources using terraform on Azure Kubernetes Service by applying the following commands:

```
terraform init  
terraform plan  
terraform apply
```

```
deepgamit@Deeps-MacBook-Pro AKS % terraform plan -var-file="azure.tfvars"  
Refreshing Terraform state in-memory prior to plan...  
The refreshed state will be used to calculate this plan, but will not be  
persisted to local or remote state storage.  
  
An execution plan has been generated and is shown below.  
Resource actions are indicated with the following symbols:  
+ create  
  
Terraform will perform the following actions:  
  
# azurerm_kubernetes_cluster.example will be created  
+ resource "azurerm_kubernetes_cluster" "example" {  
  + dns_prefix = "finalproject-k8s"  
  + fqdn  
  + id  
  + kube_admin_config  
  + kube_admin_config_raw  
  + kube_config  
  + kube_config_raw  
  + kubernetes_version  
  + location  
  + name  
  + node_resource_group  
  + portal_fqdn  
  + private_cluster_enabled  
  + private_cluster_public_fqdn_enabled = false  
  + private_dns_zone_id  
  + private_fqdn  
  + private_link_enabled  
  + resource_group_name  
  + sku_tier  
  + tags  
    + "Environment" = "Production"  
}  
  
+ addon_profile {  
  + aci_connector_linux {  
    + enabled = (known after apply)  
    + subnet_name = (known after apply)  
  }  
  
  + azure_keyvault_secrets_provider {  
    + enabled = (known after apply)  
    + secret_identity = (known after apply)  
    + secret_rotation_enabled = (known after apply)  
    + secret_rotation_interval = (known after apply)  
  }  
  
  + azure_policy {
```

```
    + outbound_ports_allocated = (known after apply)  
  }  
  
  + nat_gateway_profile {  
    + effective_outbound_ips = (known after apply)  
    + idle_timeout_in_minutes = (known after apply)  
    + managed_outbound_ip_count = (known after apply)  
  }  
  
  + role_based_access_control {  
    + enabled = (known after apply)  
  }  
  
  + azure_active_directory {  
    + admin_group_object_ids = (known after apply)  
    + azure_rbac_enabled = (known after apply)  
    + client_app_id = (known after apply)  
    + managed = (known after apply)  
    + server_app_id = (known after apply)  
    + server_app_secret = (sensitive value)  
    + tenant_id = (known after apply)  
  }  
  
  + service_principal {  
    + client_id = "89db72c0-b039-42f3-a521-0910ce0459ca"  
    + client_secret = (sensitive value)  
  }  
  
  + windows_profile {  
    + admin_password = (sensitive value)  
    + admin_username = (known after apply)  
    + license = (known after apply)  
  }  
  
# azurerm_resource_group.example will be created  
+ resource "azurerm_resource_group" "example" {  
  + id = (known after apply)  
  + location = "eastus"  
  + name = "finalproject-k8s-resources"  
}  
  
Plan: 2 to add, 0 to change, 0 to destroy.  
  
Changes to Outputs:  
+ client_certificate = (known after apply)  
+ client_key = (known after apply)  
+ cluster_ca_certificate = (known after apply)  
+ host = (known after apply)  
+ id = (known after apply)  
+ kube_config = (sensitive value)
```


Step 3: Verify that the worker nodes are joining the cluster:

```
kubectl get nodes
```

```
deepgamt@Deeps-MacBook-Pro AKS % kubectl get nodes
NAME           STATUS  ROLES   AGE    VERSION
aks-default-24425734-vmss000000  Ready   agent   14m   v1.20.9
deepgamt@Deeps-MacBook-Pro AKS %
```

Check Resource utilisation:

```
kubectl top nodes
```

```
deepgamt@Deeps-MacBook-Pro AKS % kubectl top nodes
NAME           CPU(cores)   CPU%   MEMORY(bytes)  MEMORY%
aks-default-24425734-vmss000000  200m        10%      773Mi       16%
deepgamt@Deeps-MacBook-Pro AKS %
```

Get all deployed pods:

```
kubectl get deploy --all-namespaces
```

```
deepgamt@Deeps-MacBook-Pro AKS % kubectl get deploy --all-namespaces
NAMESPACE     NAME        READY   UP-TO-DATE   AVAILABLE   AGE
kube-system   coredns     2/2     2            2           19m
kube-system   coredns-autoscaler 1/1     1            1           19m
kube-system   metrics-server 1/1     1            1           19m
kube-system   tunnelfront   1/1     1            1           19m
deepgamt@Deeps-MacBook-Pro AKS %
```

Step 4: Now that we have our cluster up and running on AKS, we deploy our rockstar app by combining the deployment and service files and applying:

```
kubectl apply -f rockstar-react-fe-combo.yaml
```

```
deepgamt@Deeps-MacBook-Pro AKS % kubectl apply -f rockstar-react-fe-combo.yaml
deployment.apps/rockstar-react-fe created
service/rockstar-react-fe-lb created
deepgamt@Deeps-MacBook-Pro AKS %
```

Step 5: Get the deployment & service details:

```
kubectl describe deploy rockstar-react-fe
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl describe deploy rockstar-react-fe

Name:           rockstar-react-fe
Namespace:      default
CreationTimestamp: Thu, 16 Dec 2021 08:58:37 -0500
Labels:          <none>
Annotations:    deployment.kubernetes.io/revision: 1
Selector:        app=rockstar-react-fe
Replicas:       1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 15
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels:       app=rockstar-react-fe
  Annotations: prometheus.io/path: /status/format/prometheus
                prometheus.io/port: 80
                prometheus.io/scrape: true
  Containers:
    rockstar-react-fe:
      Image:      deep2105/rockstar-react-fe:minikube
      Port:       80/TCP
      Host Port: 0/TCP
      Limits:
        cpu: 500m
      Requests:
        cpu: 200m
      Environment: <none>
      Mounts:      <none>
      Volumes:     <none>
  Conditions:
    Type     Status  Reason
    ----   ----   -----
    Available  True    MinimumReplicasAvailable
    Progressing  True    NewReplicaSetAvailable
  OldReplicaSets: <none>
  NewReplicaSet:  rockstar-react-fe-69ddcc89f7 (1/1 replicas created)
Events:
  Type     Reason            Age   From           Message
  ----   -----   ----   ----   -----
  Normal  ScalingReplicaSet  53s   deployment-controller  Scaled up replica set rockstar-react-fe-69ddcc89f7 to 1
```

kubectl get svc

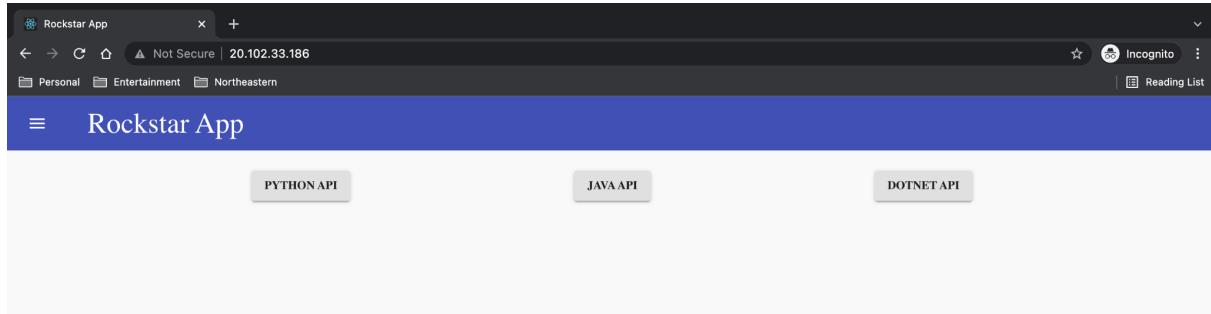
kubectl describe svc rockstar-react-fe-lb

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl get svc
NAME         TYPE      CLUSTER-IP    EXTERNAL-IP    PORT(S)      AGE
kubernetes   ClusterIP  10.0.0.1    <none>        443/TCP    30m
rockstar-react-fe-lb  LoadBalancer  10.0.10.103  20.102.33.186  80:32354/TCP  6m18s
deepgamit@Deeps-MacBook-Pro AKS %
```

deepgamit@Deeps-MacBook-Pro AKS % kubectl describe svc rockstar-react-fe-lb

```
Name:           rockstar-react-fe-lb
Namespace:      default
Labels:          <none>
Annotations:    <none>
Selector:        app=rockstar-react-fe
Type:           LoadBalancer
IP Families:   <none>
IP:             10.0.10.103
IPs:            10.0.10.103
LoadBalancer Ingress: 20.102.33.186
Port:           <unset> 80/TCP
TargetPort:     80/TCP
NodePort:       <unset> 32354/TCP
Endpoints:      10.244.0.7:80
Session Affinity: None
External Traffic Policy: Cluster
Events:
  Type     Reason            Age   From           Message
  ----   -----   ----   ----   -----
  Normal  EnsuringLoadBalancer 9m44s  service-controller  Ensuring load balancer
  Normal  EnsuredLoadBalancer  9m27s  service-controller  Ensured load balancer
```

APPLICATION RUNNING ON IP: 20.102.33.186



Step 6: After deploying the application successfully, let's add horizontal autoscaling by creating a horizontal pod autoscaler file.

```
io.k8s.api.autoscaling.v1.HorizontalPodAutoscaler (v1@horizontalpodautoscaler.json)
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: rockstar-react
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: rockstar-react-fe
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 5
```

Now create the auto scaler by:

```
kubectl create -f rockstar-app-scaler.yaml
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl create -f rockstar-app-scaler.yaml
horizontalpodautoscaler.autoscaling/rockstar-react created
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 7: Now lets test if auto scaling works by running the following:

```
export SERVICE_IP=$(kubectl get svc rockstar-react-fe-lb
-o json | jq -r '.status.loadBalancer.ingress[].ip')
```

Here we export the service ip as an environment variable. Now we can use this variable for load testing with the following command:

```
ab -c 100 -n 100000 http://${SERVICE_IP}/
```

```
deepgamit@Deeps-MacBook-Pro AKS % ab -c 100 -n 100000 http://${SERVICE_IP}/
```

```
This is ApacheBench, Version 2.3 <$Revision: 1879490 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
```

```
Benchmarking 20.102.33.186 (be patient)
Completed 10000 requests
Completed 20000 requests
[
```

Step 8: Check the stats with:

```
kubectl get hpa rockstar-react
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl get hpa rockstar-react
NAME          REFERENCE           TARGETS      MINPODS     MAXPODS     REPLICAS   AGE
rockstar-react Deployment/rockstar-react-fe  0%/5%       1           10          1          2m42s
deepgamit@Deeps-MacBook-Pro AKS % [
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl get hpa
NAME          REFERENCE           TARGETS      MINPODS     MAXPODS     REPLICAS   AGE
rockstar-react Deployment/rockstar-react-fe  13%/5%      1           10          3          55m
deepgamit@Deeps-MacBook-Pro AKS % [
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl get hpa
NAME          REFERENCE           TARGETS      MINPODS     MAXPODS     REPLICAS   AGE
rockstar-react Deployment/rockstar-react-fe  4%/5%       1           10          3          56m
deepgamit@Deeps-MacBook-Pro AKS % kubectl get hpa
NAME          REFERENCE           TARGETS      MINPODS     MAXPODS     REPLICAS   AGE
rockstar-react Deployment/rockstar-react-fe  5%/5%       1           10          3          57m
deepgamit@Deeps-MacBook-Pro AKS % kubectl get hpa
NAME          REFERENCE           TARGETS      MINPODS     MAXPODS     REPLICAS   AGE
rockstar-react Deployment/rockstar-react-fe  5%/5%       1           10          3          57m
deepgamit@Deeps-MacBook-Pro AKS % kubectl get hpa
NAME          REFERENCE           TARGETS      MINPODS     MAXPODS     REPLICAS   AGE
rockstar-react Deployment/rockstar-react-fe  5%/5%       1           10          3          58m
deepgamit@Deeps-MacBook-Pro AKS % [
```

Step 9: Now we are ready to deploy prometheus for monitoring. Let's create a different namespace to separate monitoring pods from k8s app pods.

```
kubectl create namespace monitoring
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl create namespace monitoring
namespace/monitoring created
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 10: For Prometheus to fetch metrics from Kubernetes' APIs, we will assign the cluster reader permission. For that, we create a cluster role:

```
v1@clusterrolebinding.json | v1@clusterrolebinding.json
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: prometheus
rules:
- apiGroups: [""]
  resources:
  - nodes
  - nodes/proxy
  - services
  - endpoints
  - pods
  verbs: ["get", "list", "watch"]
- apiGroups:
  - extensions
  resources:
  - ingresses
  verbs: ["get", "list", "watch"]
- nonResourceURLs: ["/metrics"]
  verbs: ["get"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: prometheus
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: prometheus
subjects:
- kind: ServiceAccount
  name: default
  namespace: monitoring
```

Create permissions with the following command:

```
kubectl create -f clusterRole.yaml
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl create -f clusterRole.yaml
clusterrole.rbac.authorization.k8s.io/prometheus created
clusterrolebinding.rbac.authorization.k8s.io/prometheus created
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 11: After giving proper permissions, we create a config map for our cluster with prometheus scrape config which will be mounted to Prometheus container in /etc/prometheus as files `prometheus.yml` and `prometheus.rules`

```
io.k8s.api.core.v1.ConfigMap (v1@configmap.json)
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-server-conf
  labels:
    name: prometheus-server-conf
  namespace: monitoring
data:
  prometheus.rules: |-  
    groups:  
      - name: devopscube demo alert  
        rules:  
          - alert: High Pod Memory  
            expr: sum(container_memory_usage_bytes) > 1  
            for: 1m  
            labels:  
              severity: slack  
            annotations:  
              summary: High Memory Usage
  prometheus.yml: |-  
    global:  
      scrape_interval: 5s  
      evaluation_interval: 5s  
    rule_files:  
      - /etc/prometheus/prometheus.rules  
    alerting:  
      alertmanagers:  
        - scheme: http  
        static_configs:  
          - targets:  
            - "alertmanager.monitoring.svc:9093"  
  
    scrape_configs:  
      - job_name: 'kubernetes-apiservers'  
  
        kubernetes_sd_configs:  
          - role: endpoints  
        scheme: https  
  
        tls_config:  
          ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt  
          bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
```

Create config map with the following command:

```
kubectl create -f config-map.yaml
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl create -f config-map.yaml  
configmap/prometheus-server-conf created  
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 12: Create Prometheus deployment after creating the config map

```
io.k8s.api.apps.v1.Deployment (v1@deployment.json)
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus-server
  namespace: monitoring
spec:
  selector:
    matchLabels:
      app: prometheus-server
  replicas: 1
  template:
    metadata:
      labels:
        app: prometheus-server
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus:v2.12.0
          args:
            - "--config.file=/etc/prometheus/prometheus.yml"
            - "--storage.tsdb.path=/prometheus/"
          ports:
            - containerPort: 9090
          volumeMounts:
            - name: prometheus-config-volume
              mountPath: /etc/prometheus/
            - name: prometheus-storage-volume
              mountPath: /prometheus/
      volumes:
        - name: prometheus-config-volume
          configMap:
            defaultMode: 420
            name: prometheus-server-conf

        - name: prometheus-storage-volume
          emptyDir: {}
```

Deploy on monitoring namespace using prometheus-deployment file:

```
kubectl create -f prometheus-deployment.yaml
--namespace=monitoring
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl create -f prometheus-deployment.yaml --namespace=monitoring
deployment.apps/prometheus-server created
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 13: Checking all the deployments

```
kubectl get deployments --all-namespaces
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl get deployments --all-namespaces
NAMESPACE      NAME            READY   UP-TO-DATE   AVAILABLE   AGE
default        rockstar-react-fe   1/1     1           1           38m
kube-system    coredns          2/2     2           2           61m
kube-system    coredns-autoscaler 1/1     1           1           61m
kube-system    metrics-server    1/1     1           1           61m
kube-system    tunnelfront       1/1     1           1           61m
monitoring     prometheus-server 1/1     1           1           91s
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 14: Let's expose Prometheus as a service so that we can access the Prometheus dashboard over an IP

```
io.k8s.api.core.v1.Service (v1@service.json)
apiVersion: v1
kind: Service
metadata:
  name: prometheus-service
  namespace: monitoring
  annotations:
    | prometheus.io/scrape: 'true'
    | prometheus.io/port:    '9090'
spec:
  selector:
    | app: prometheus-server
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 9090
```

Create service endpoint with the following command:

```
kubectl create -f prometheus-service.yaml
--namespace=monitoring
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl create -f prometheus-service.yaml --namespace=monitoring
service/prometheus-service created
deepgamit@Deeps-MacBook-Pro AKS %
```

Step 15: Get the IP of Prometheus service with:

```
kubectl get svc --namespace=monitoring
```

```
deepgamit@Deeps-MacBook-Pro AKS % kubectl get svc --namespace=monitoring
NAME           TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)        AGE
prometheus-service   LoadBalancer  10.0.158.166  52.151.232.12  80:31301/TCP  102s
deepgamit@Deeps-MacBook-Pro AKS %
```

PROMETHEUS DASHBOARD IP: 52.151.232.12

Rockstar App × Prometheus Time Series Collector × +

← → C ⌘ ⌘ Not Secure | 52.151.232.12/graph

Personal Entertainment Northeastern

Incognito Reading List

Prometheus Alerts Graph Status Help

Enable query history

Expression (press Shift+Enter for newlines)

Execute - insert metric at cursor -

Graph Console

Moment

Element Value

no data

Add Graph Remove Graph

TARGETS

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
https://kubernetes.default.svc:443/api/v1/nodes/aks-default-24425734-vmss000000/proxy/metrics/cadvisor	UP	agentpool="default" beta_kubernetes_io_arch="amd64" beta_kubernetes_io_instance_type="Standard_D2_v2" beta_kubernetes_io_os="linux" failure_domain_beta_kubernetes_io_regions="eastus" failure_domain_beta_kubernetes_io_zones="0" instances="aks-default-24425734-vmss000000" job="kubernetes-cadvisor" kubernetes_azure_com_agentpool="default" kubernetes_azure_com_cluster="MC_finalproject-k8s-resources_finalproject-k8s_eastus" kubernetes_azure_com_mode="system" kubernetes_azure_com_node_image_versions="AKSUbuntu-1804containerd-2021.11.27" kubernetes_azure_com_os_sku="Ubuntu" kubernetes_azure_com_role="agent" kubernetes_azure_com_storageprofile="managed" kubernetes_azure_com_storagelayer="Standard_LRS" kubernetes_io_arch="amd64" kubernetes_io_hostname="aks-default-24425734-vmss000000" kubernetes_io_os="linux" kubernetes_io_role="agent" node_kubernetes_io_instance_type="Standard_D2_v2" storageprofile="managed" storagelayer="Standard_LRS" topology_kubernetes_io_region="eastus" topology_kubernetes_io_zone="0"	3.882s ago	49.06ms	
https://kubernetes.default.svc:443/api/v1/nodes/aks-default-24425734-vmss000000/proxy/metrics	UP	agentpool="default" beta_kubernetes_io_arch="amd64" beta_kubernetes_io_instance_type="Standard_D2_v2" beta_kubernetes_io_os="linux" failure_domain_beta_kubernetes_io_regions="eastus" failure_domain_beta_kubernetes_io_zones="0"	4.283s ago	61ms	

SEE ALL METRICS: 52.151.232.12/metrics

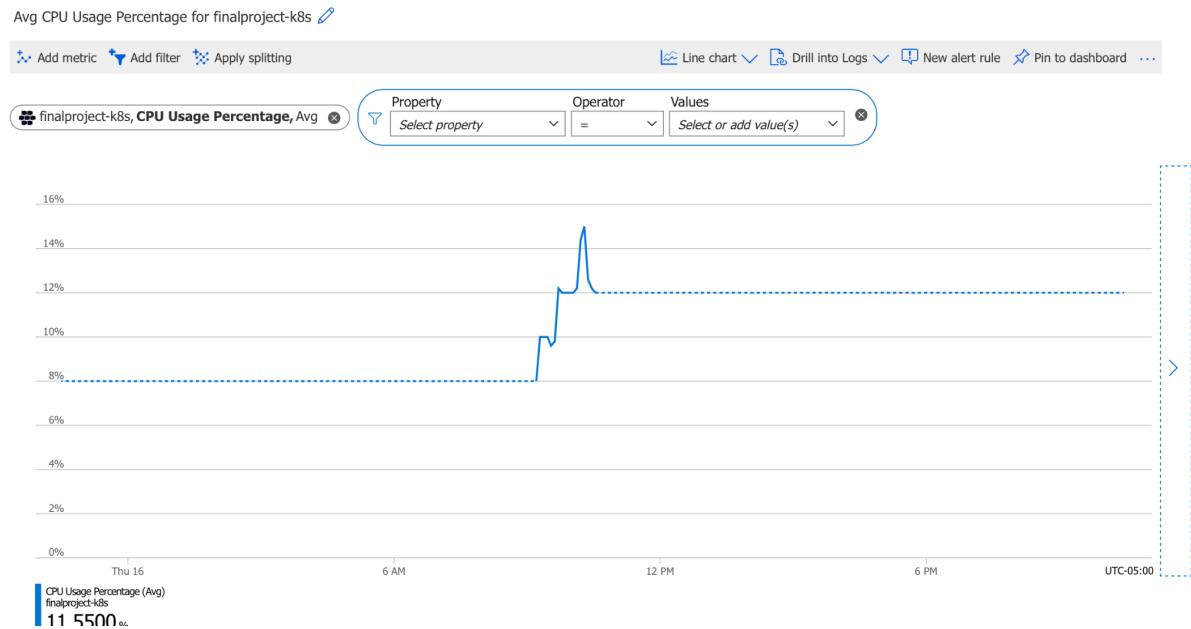
```
# HELP go_gc_duration_seconds A summary of the GC invocation durations.
# TYPE go_gc_duration_seconds summary
go_gc.duration.seconds.quantile="0.05" 9.9e-06
go_gc.duration.seconds.quantile="0.25" 1.9901e-05
go_gc.duration.seconds.quantile="0.5" 3.5202e-05
go_gc.duration.seconds.quantile="0.75" 5.1802e-05
go_gc.duration.seconds.sum 0.002230097
go_gc.duration.seconds.count 34
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go.goroutines 1
# HELP go_info Information about the Go environment.
# TYPE go_info gauge
go.info(version="go1.12.8") 1
# HELP go_memstats_alloc_bytes Number of bytes allocated and still in use.
# TYPE go_memstats_alloc_bytes gauge
go.memstats.alloc.bytes 3.8475648e+07
# HELP go_memstats_alloc_bytes_total Total number of bytes allocated, even if freed.
# TYPE go_memstats_alloc_bytes_total counter
go.memstats.alloc.bytes.total 5.18216e+08
# HELP go_memstats_buck_hash_table_bytes Number of bytes used by the profiling bucket hash table.
# TYPE go_memstats_buck_hash_table_bytes gauge
go.memstats.buck.hash.table.bytes 1.540041e+06
# HELP go_memstats_frees_total Total number of frees.
# TYPE go_memstats_frees_total counter
go.memstats.frees.total 3.117768e+06
# HELP go_memstats_gc_cpu_fraction The fraction of this program's available CPU time used by the GC since the program started.
# TYPE go_memstats_gc_cpu_fraction gauge
go.memstats.gc.cpu.fraction 2.976092793290418e-05
# HELP go_memstats_gc_sys_bytes Number of bytes used for garbage collection system metadata.
# TYPE go_memstats_gc_sys_bytes gauge
go.memstats.gc.sys.bytes 2.625776e+06
# HELP go_memstats_heap_alloc_bytes Number of heap bytes allocated and still in use.
# TYPE go_memstats_heap_alloc_bytes gauge
go.memstats.heap.alloc.bytes 3.8475648e+07
# HELP go_memstats_heap_idle_bytes Number of heap bytes waiting to be used.
# TYPE go_memstats_heap_idle_bytes gauge
go.memstats.heap.idle.bytes 2.3986176e+07
# HELP go_memstats_heap_inuse_bytes Number of heap bytes that are in use.
# TYPE go_memstats_heap_inuse_bytes gauge
go.memstats.heap.inuse.bytes 2.041344e+07
# HELP go_memstats_heap_objects Number of allocated objects.
# TYPE go_memstats_heap_objects gauge
go.memstats.heap.objects 249938
# HELP go_memstats_heap_released_bytes Number of heap bytes released to OS.
# TYPE go_memstats_heap_released_bytes gauge
go.memstats.heap.released.bytes 0
# HELP go_memstats_heap_sys_bytes Number of heap bytes obtained from system.
# TYPE go_memstats_heap_sys_bytes gauge
go.memstats.heap.sys.bytes 6.602752e+07
# HELP go_memstats_last_gc_time_seconds Number of seconds since 1970 of last garbage collection.
```

METRICS OF INTEREST ON PROMETHEUS

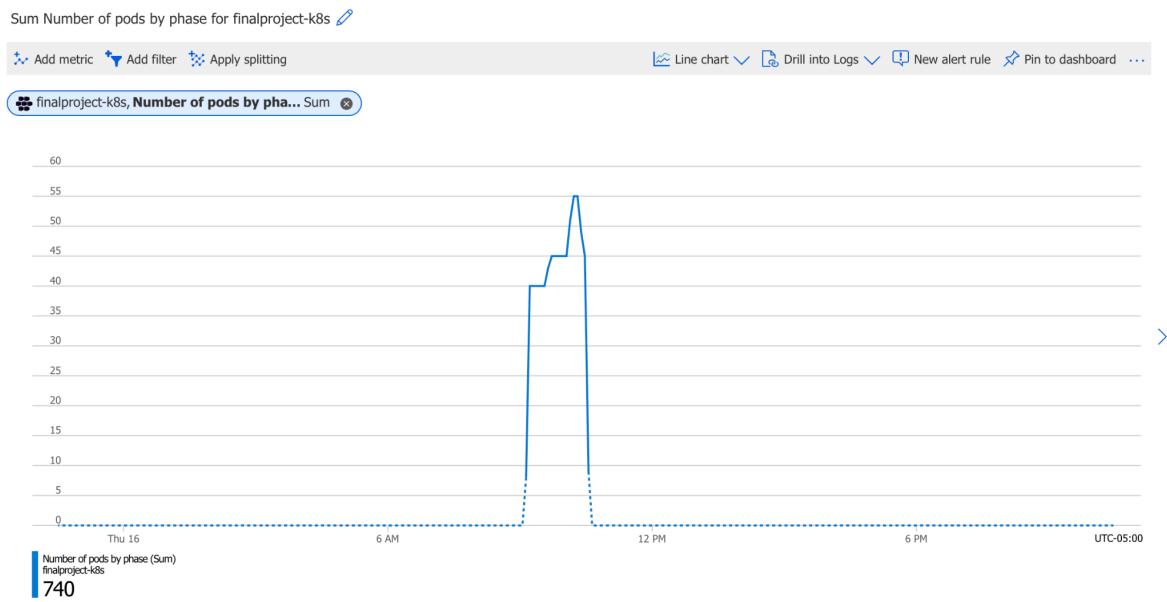


METRICS ON AZURE

Average CPU Usage Percentage



Sum Of Number of Pods in Phase



Deploying the Rockstar App to EKS cluster with CloudWatch for Monitoring

Step 1: We deploy the resources using terraform on AWS Elastic Kubernetes Service by applying the following commands:

```
terraform init  
terraform plan  
terraform apply
```

```
+ self          = false  
+ source_security_group_id = (known after apply)  
+ to_port       = 443  
+ type         = "ingress"  
  
# aws_security_group_rule.final-node-ingress-cluster will be created  
+ resource "aws_security_group_rule" "final-node-ingress-cluster" {  
  + description      = "Allow worker Kubelets and pods to receive communication from the worker nodes"  
  + from_port        = 1025  
  + id               = (known after apply)  
  + protocol         = "tcp"  
  + security_group_id = (known after apply)  
  + self              = false  
  + source_security_group_id = (known after apply)  
  + to_port          = 65535  
  + type              = "ingress"  
  
# aws_security_group_rule.final-node-ingress-self will be created  
+ resource "aws_security_group_rule" "final-node-ingress-self" {  
  + description      = "Allow node to communicate with each other"  
  + from_port        = 0  
  + id               = (known after apply)  
  + protocol         = "-1"  
  + security_group_id = (known after apply)  
  + self              = false  
  + source_security_group_id = (known after apply)  
  + to_port          = 65535  
  + type              = "ingress"  
  
Plan: 21 to add, 0 to change, 0 to destroy.  
  
Changes to Outputs:  
  config_map_aws_auth = (known after apply)  
  kubeconfig          = (known after apply)  
  kubeverSION         = (known after apply)
```

```
aws_eks_cluster.final: Still creating... [7m10s elapsed]  
aws_eks_cluster.final: Still creating... [7m20s elapsed]  
aws_eks_cluster.final: Still creating... [7m30s elapsed]  
aws_eks_cluster.final: Still creating... [7m40s elapsed]  
aws_eks_cluster.final: Still creating... [7m50s elapsed]  
aws_eks_cluster.final: Still creating... [8m0s elapsed]  
aws_eks_cluster.final: Still creating... [8m10s elapsed]  
aws_eks_cluster.final: Still creating... [8m20s elapsed]  
aws_eks_cluster.final: Still creating... [8m30s elapsed]  
aws_eks_cluster.final: Still creating... [8m40s elapsed]  
aws_eks_cluster.final: Still creating... [8m50s elapsed]  
aws_eks_cluster.final: Still creating... [9m0s elapsed]  
aws_eks_cluster.final: Still creating... [9m10s elapsed]  
aws_eks_cluster.final: Still creating... [9m20s elapsed]  
aws_eks_cluster.final: Still creating... [9m30s elapsed]  
aws_eks_cluster.final: Still creating... [9m40s elapsed]  
aws_eks_cluster.final: Still creating... [9m50s elapsed]  
aws_eks_cluster.final: Still creating... [10m0s elapsed]  
aws_eks_cluster.final: Creation complete after 10m6s [id=terraform-data.aws_ami.eks-worker: Reading...]  
data.aws_ami.eks-worker: Read complete after 0s [id=ami-019904275e]  
aws_launch_configuration.final: Creating...  
aws_launch_configuration.final: Creation complete after 2s [id=terraform-data.aws_launch_configuration.eks-worker: Reading...]  
aws_autoscaling_group.final: Creating...  
aws_autoscaling_group.final: Still creating... [10s elapsed]  
aws_autoscaling_group.final: Still creating... [20s elapsed]  
aws_autoscaling_group.final: Creation complete after 28s [id=terraform-data.aws_autoscaling_group.eks-worker: Reading...]  
  
Apply complete! Resources: 21 added, 0 changed, 0 destroyed.
```

Step 2: Configure kubectl locally in order to communicate with AWS Elastic Kubernetes Server. Running **terraform output kubeconfig** to generate a sample kubectl configuration to connect to our cluster

Running the following commands:

```
terraform output kubeconfig > ~/.kube/config-terraform-eks-demo  
cp ~/.kube/config-terraform-eks-demo ~/.kube/config
```

```

Outputs:
config_map_aws_auth =
apiVersion: v1
kind: ConfigMap
metadata:
  name: aws-auth
  namespace: kube-system
data:
mapRoles: |
  - rolearn: arn:aws:iam::658346791705:role/terraform-eks-final-node
    username: system:node:{{EC2PrivateDNSName}}
    groups:
      - system:bootstrappers
      - system:nodes
kubeconfig =
apiVersion: v1
clusters:
- cluster:
  server: https://5FC016FF626CF20BF7E6E8C1B91F289B.gr7.us-east-1.eks.amazonaws.com
  certificate-authority-data: LS0tLS1CRUdjTlBDRVJUSUZjQ0FURS0tLS0tCk1JSU1ekNDQWMrZ0F3SUJBZ0lCQURBTkJna3Foa2lHOXcwQkFRc0ZBREfWTJNjd0VRWURWUVEFRdwcmRKSxmwKY201bGRHvnpoN0jRYRFRJeE1USXh0akUzTURjME5Gb1hEVE14TVRje5ERTNNRGWwTkZvd0ZURVRN0kVHQTFVR0pBeE1LYTNwNa/pY5nVawNFjsY3pDQ0FTSXdeUVlKS29asSwl2Y5B0UVC0lFBRGdnRVBBRENDDQFvQ2dnRjUZBTFVBcns5nUnLQdHNSFPqUdVRST02NmFY0Elvb2REUUQ4Q1LRRZXBNB0dSTK5pUgh0MGQ4MXdzY1ZnTnMvSk1LTTj0TQkjdikNko1025XV1NBmJvWVFhJYVZYMwsdmpMZFHhcncNFTWnrUjN1cnd0VytDZXLzM3hSY2R0U21Jdja4eJ4b0puHdBTXkvccyrlNaUGFYekqrB3gzWfp5Q0FyNfpCo3dZTp4bkzsbtNrTlpJwnB1Rw@MEJCRWF1N2RabFhjCmFhGh20jQ3d6tOrGd35m1VSWRYcLjQTHew3BjUTA1aTNhjM4U19Qemz0a3ZZNeckv1MrcmZ5WwhodmhGVLIKMTFxvUvXU21rSLRVSNh6Cxh0YQ3aHhYQlpKdVxSURJUkjtyXNWgprVjdMMWRYMWVKTWUW901kwcGRDwppqMkNOV29tY5p0h1R5bk1WeXpF00F3RUFByU5D0TUVBd0RnWURWUjB0QVFl0JBUURBZ0tTUE4R0ExWVRfd0VCCi93UZNUQ1CQWY4d0hRWURWUjBP0kjZRUZYVMybxnMNTVpVnBGERMz0m5xdzBuWLnhWkhN0TBHQ1LNxR1NjYjMKRFFFQkN3VUFBNELCQVFBynZfaFlzTLbjMGhMvnr5Y3ZRNGFwagdjRDNVVEhuwFnqUU9yZQRQwVLGeVFcdmc5NgptSw0102phvZ50TYzRGVvadBu21RNFFid25NNW83T0JL0TDXNvJCD2J1OFBhwDZKbwh0dg02emFDTX1DTypxChhNem9ya201Rzg3U0pmNLFSVHZNN01LcLNyUnJrUH1Qm5L29NdE1rL3p0UVJKd2FqRXFM0zJrZUdLRkF0MEYKNUVuTlZTbdBz0UhwsTR01L2hqdmJ1VDBJQ045Z2g1bDMSNTZQk2NreVz0d21YU0k5Zm16cExUbGRXYUpUtEn2awo3WnZHSVJjcDRmMm5SdU0ycnhJeWn4S3dKMFhU2FxRkNVmZRSj09KK003RFV5bm1wT1JRT0paNFTYUpudm3CdjUn1ZU5VWVNsUDVUeVSUihR0tnRzNEVHFMTRLVDJvLwotLS0tLUVORCBDRVJUSUZjQ0FURS0tLS0tCg==

name: kubernetes
contexts:
- context:
  cluster: kubernetes
  user: aws
  name: aws
current-context: aws
kind: Config
preferences: {}
users:
- name: aws
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      command: aws-iam-authenticator
      args:
        - "token"
        - "-i"

```

Step 3: Verify that the worker nodes are joining the cluster:

```
kubectl get nodes
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl get nodes
NAME                  STATUS   ROLES      AGE      VERSION
ip-10-0-0-216.ec2.internal   Ready    <none>    66s     v1.21.5-eks-bc4871b
ip-10-0-1-176.ec2.internal   Ready    <none>    65s     v1.21.5-eks-bc4871b
deepgamit@Deeps-MacBook-Pro EKS %
```

Get all deployed pods:

```
kubectl get deploy --all-namespaces
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl get deploy --all-namespaces
NAMESPACE      NAME      READY   UP-TO-DATE   AVAILABLE   AGE
kube-system   coredns   2/2     2           2           9m42s
deepgamit@Deeps-MacBook-Pro EKS %
```

Step 4: Now that we have our cluster up and running on EKS, we deploy our rockstar app by combining the deployment and service files and applying:

```
kubectl apply -f rockstar-react-fe-combo.yaml
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl apply -f rockstar-react-fe-combo.yaml
deployment.apps/rockstar-react-fe created
service/rockstar-react-fe-lb created
deepgamit@Deeps-MacBook-Pro EKS %
```

Step 5: Get service details:

```
kubectl get svc
kubectl get pods
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl get svc
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP
PORT(S)        AGE
kubernetes     ClusterIP   172.20.0.1    <none>
  443/TCP       11m
rockstar-react-fe-lb   LoadBalancer   172.20.181.163  a3530826f509a499ea27c6fce154b5b-182992700.us-east-1.elb.amazonaws.com  80:32265/TCP  43s
deepgamit@Deeps-MacBook-Pro EKS %
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
rockstar-react-fe-76bd554599-47rkc   1/1     Running   0          2m14s
rockstar-react-fe-76bd554599-4b85c   1/1     Running   0          2m14s
rockstar-react-fe-76bd554599-cd8fk   1/1     Running   0          2m14s
deepgamit@Deeps-MacBook-Pro EKS %
```

APPLICATION RUNNING ON IP:

a3530826f509a499ea27c6fce154b5b-182992700.us-east-1.elb.amazonaws.com

Setup CloudWatch Agent to collect metrics

Step 6: Create a namespace for CloudWatch:

```
kubectl apply -f cwagent-namespace.yaml  
kubectl get namespaces
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl apply -f cwagent-namespace.yaml  
namespace/amazon-cloudwatch created  
deepgamit@Deeps-MacBook-Pro EKS % kubectl get namespaces  
NAME STATUS AGE  
amazon-cloudwatch Active 9s  
default Active 16m  
kube-node-lease Active 16m  
kube-public Active 16m  
kube-system Active 16m  
deepgamit@Deeps-MacBook-Pro EKS % █
```

Step 7: Create a service account in the cluster with ClusterRole & ClusterRoleBinding

```
kubectl apply -f cwagent-serviceaccount.yaml
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl apply -f cwagent-serviceaccount.yaml  
serviceaccount/cloudwatch-agent created  
clusterrole.rbac.authorization.k8s.io/cloudwatch-agent-role created  
clusterrolebinding.rbac.authorization.k8s.io/cloudwatch-agent-role-binding created  
deepgamit@Deeps-MacBook-Pro EKS % █
```

```
# create cwagent service account and role binding  
apiVersion: v1  
kind: ServiceAccount  
metadata:  
  name: cloudwatch-agent  
  namespace: amazon-cloudwatch  
  
---  
kind: ClusterRole  
apiVersion: rbac.authorization.k8s.io/v1  
metadata:  
  name: cloudwatch-agent-role  
rules:  
  - apiGroups: [""]  
    resources: ["pods", "nodes", "endpoints"]  
    verbs: ["list", "watch"]  
  - apiGroups: ["apps"]  
    resources: ["replicasets"]  
    verbs: ["list", "watch"]  
  - apiGroups: ["batch"]  
    resources: ["jobs"]  
    verbs: ["list", "watch"]  
  - apiGroups: [""]  
    resources: ["nodes/proxy"]  
    verbs: ["get"]  
  - apiGroups: [""]  
    resources: ["nodes/stats", "configmaps", "events"]  
    verbs: ["create"]  
  - apiGroups: [""]  
    resources: ["configmaps"]  
    resourceNames: ["cwagent-clusterleader"]  
    verbs: ["get", "update"]  
  
---  
kind: ClusterRoleBinding  
apiVersion: rbac.authorization.k8s.io/v1  
metadata:  
  name: cloudwatch-agent-role-binding  
subjects:  
  - kind: ServiceAccount  
    name: cloudwatch-agent  
    namespace: amazon-cloudwatch  
roleRef:  
  kind: ClusterRole  
  name: cloudwatch-agent-role  
  apiGroup: rbac.authorization.k8s.io
```

Step 8: Create a ConfigMap for the CloudWatch agent

```
kubectl apply -f cwagent-configmap.yaml
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl apply -f cwagent-configmap.yaml
configmap/cwagentconfig created
deepgamit@Deeps-MacBook-Pro EKS %
```

```
# create configmap for cwagent config
apiVersion: v1
data:
  # Configuration is in Json format. No matter what configure change you make,
  # please keep the Json blob valid.
  cwagentconfig.json: |
    {
      "logs": {
        "metrics_collected": {
          "kubernetes": {
            "cluster_name": "terraform-eks-final",
            "metrics_collection_interval": 60
          }
        },
        "force_flush_interval": 5
      }
    }
kind: ConfigMap
metadata:
  name: cwagentconfig
  namespace: amazon-cloudwatch
```

Step 9: Deploy the CloudWatch agent as a DaemonSet

```
kubectl apply -f cwagent-daemonset.yaml
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl apply -f cwagent-daemonset.yaml
daemonset.apps/cloudwatch-agent created
deepgamit@Deeps-MacBook-Pro EKS %
```

```
io.k8s.api.apps.v1.DaemonSet (v1@daemonset.json)
# deploy cwagent as daemonset
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: cloudwatch-agent
  namespace: amazon-cloudwatch
spec:
  selector:
    matchLabels:
      name: cloudwatch-agent
  template:
    metadata:
      labels:
        name: cloudwatch-agent
    spec:
      containers:
        - name: cloudwatch-agent
          image: amazon/cloudwatch-agent:1.247348.0b251302
          #ports:
          #  - containerPort: 8125
          #    hostPort: 8125
          #    protocol: UDP
          resources:
            limits:
              cpu: 200m
              memory: 200Mi
            requests:
              cpu: 200m
              memory: 200Mi
          # Please don't change below envs
          env:
            - name: HOST_IP
              valueFrom:
                fieldRef:
                  fieldPath: status.hostIP
            - name: HOST_NAME
              valueFrom:
                fieldRef:
```

```
                - name: docker.sock
                  mountPath: /var/run/docker.sock
                  readOnly: true
                - name: varlibdocker
                  mountPath: /var/lib/docker
                  readOnly: true
                - name: containerdsock
                  mountPath: /run/containerd/containerd.sock
                  readOnly: true
                - name: sys
                  mountPath: /sys
                  readOnly: true
                - name: devdisk
                  mountPath: /dev/disk
                  readOnly: true
  volumes:
    - name: cwagentconfig
      configMap:
        name: cwagentconfig
    - name: rootfs
      hostPath:
        path: /
    - name: dockersock
      hostPath:
        path: /var/run/docker.sock
    - name: varlibdocker
      hostPath:
        path: /var/lib/docker
    - name: containerdsock
      hostPath:
        path: /run/containerd/containerd.sock
    - name: sys
      hostPath:
        path: /sys
    - name: devdisk
      hostPath:
        path: /dev/disk/
terminationGracePeriodSeconds: 60
serviceAccountName: cloudwatch-agent
```

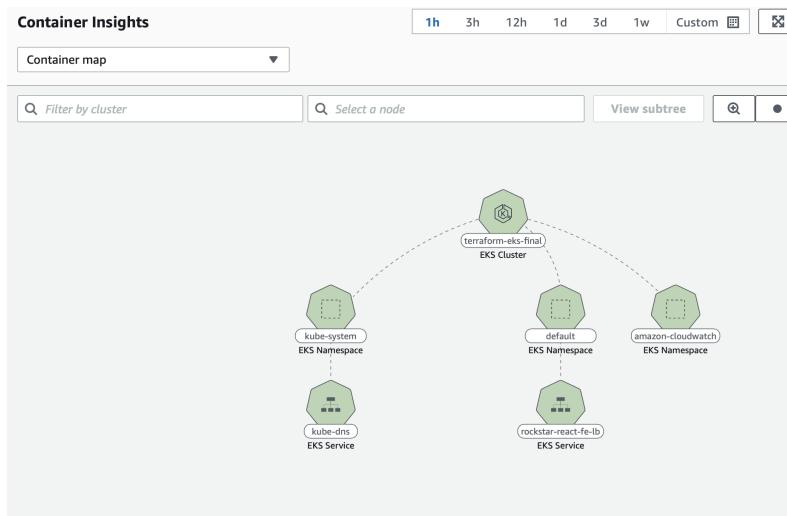
Step 10: Validate that the agent is deployed by running the following command.

```
kubectl get pods -n amazon-cloudwatch
```

```
deepgamit@Deeps-MacBook-Pro EKS % kubectl get pods -n amazon-cloudwatch
NAME                  READY   STATUS    RESTARTS   AGE
cloudwatch-agent-9sglg 1/1     Running   0          24s
cloudwatch-agent-mvnmt 1/1     Running   0          24s
deepgamit@Deeps-MacBook-Pro EKS %
```

Viewing Container Insights Metrics

Container Map

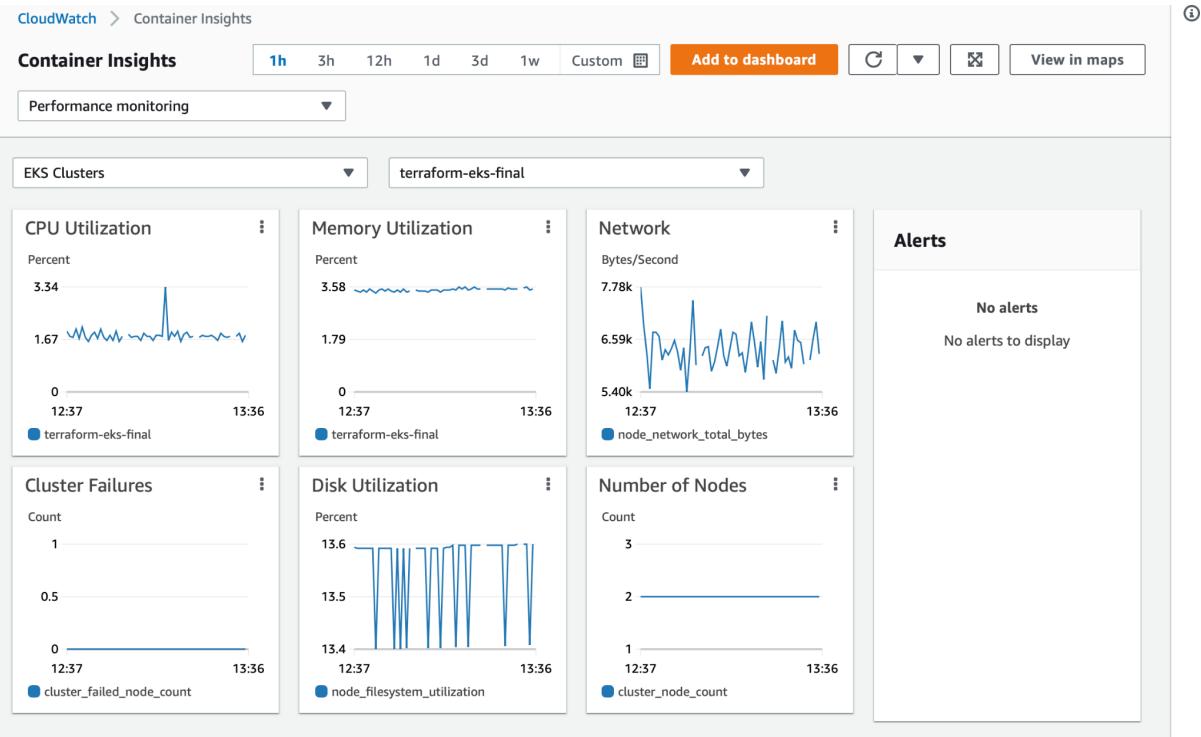


Resources

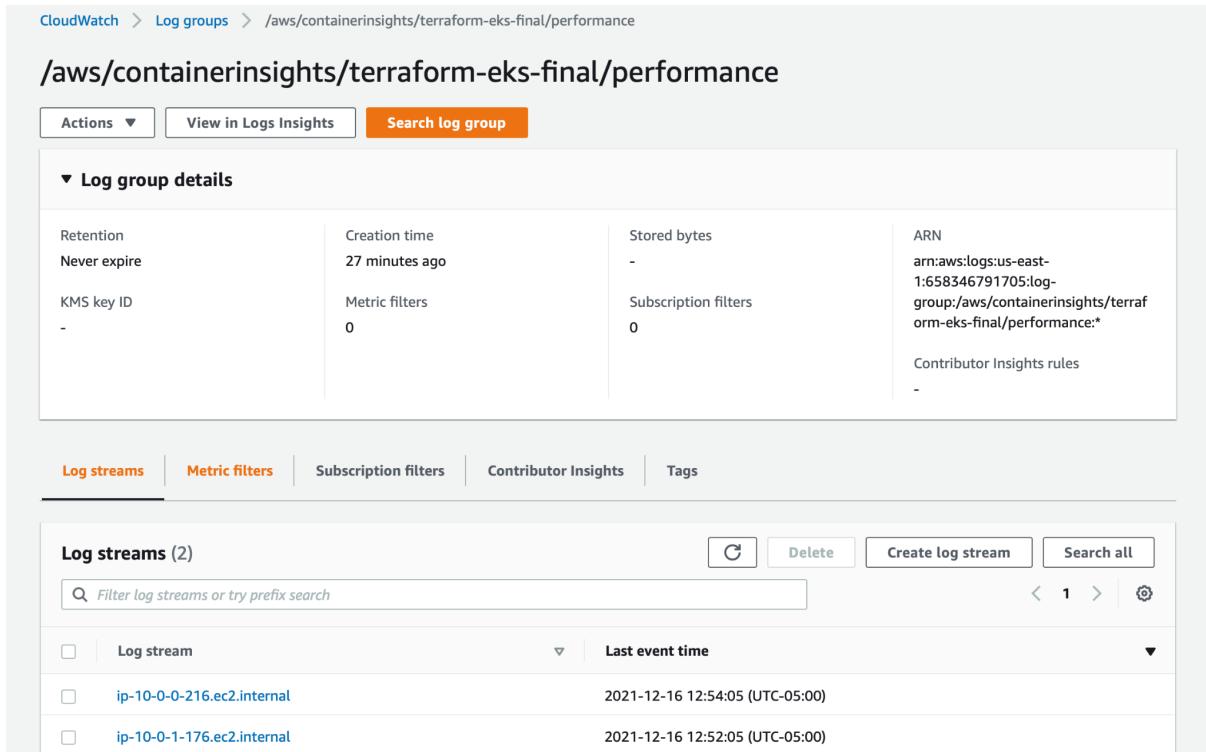
The table displays the list of resources monitored by Container Insights. It includes columns for Name, Type, Cluster name, Alarms, Prometheus, Avg CPU (%), and Avg memory (%). The data shows six entries corresponding to the nodes shown in the Container Map.

	Name	Type	Cluster name	Alarms	Prometheus	Avg CPU (%)	Avg memory (%)
○	amazon-cloudwatch	EKS Namespace	terraform-eks-final	-	-	0.1%	0.2%
○	default	EKS Namespace	terraform-eks-final	-	-	<0.1%	<0.1%
○	kube-dns	EKS Service	terraform-eks-final	-	-	<0.1%	<0.1%
○	kube-system	EKS Namespace	terraform-eks-final	-	-	<0.1%	0.1%
○	rockstar-react-fe-lb	EKS Service	terraform-eks-final	-	-	<0.1%	<0.1%
○	terraform-eks-final	EKS Cluster	terraform-eks-final	-	-	2.1%	3.5%

Performance Monitoring



Log Groups for Container Insights



Setup SNS Topic & CloudWatch Alarms

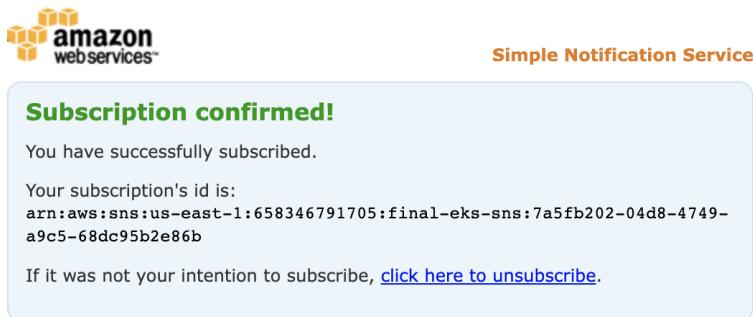
Step 1: Create a topic with name *final-eks-sns*

```
deepgamit@Deeps-MacBook-Pro EKS % aws sns create-topic --name final-eks-sns
{
    "TopicArn": "arn:aws:sns:us-east-1:658346791705:final-eks-sns"
}
deepgamit@Deeps-MacBook-Pro EKS % █
```

Step 2: Subscribe the topic to email address

```
deepgamit@Deeps-MacBook-Pro EKS % aws sns subscribe --topic-arn arn:aws:sns:us-east-1:658346791705:final-eks-sns --protocol
email --notification-endpoint deepgamit@gmail.com
{
    "SubscriptionArn": "pending confirmation"
}
deepgamit@Deeps-MacBook-Pro EKS % █
```

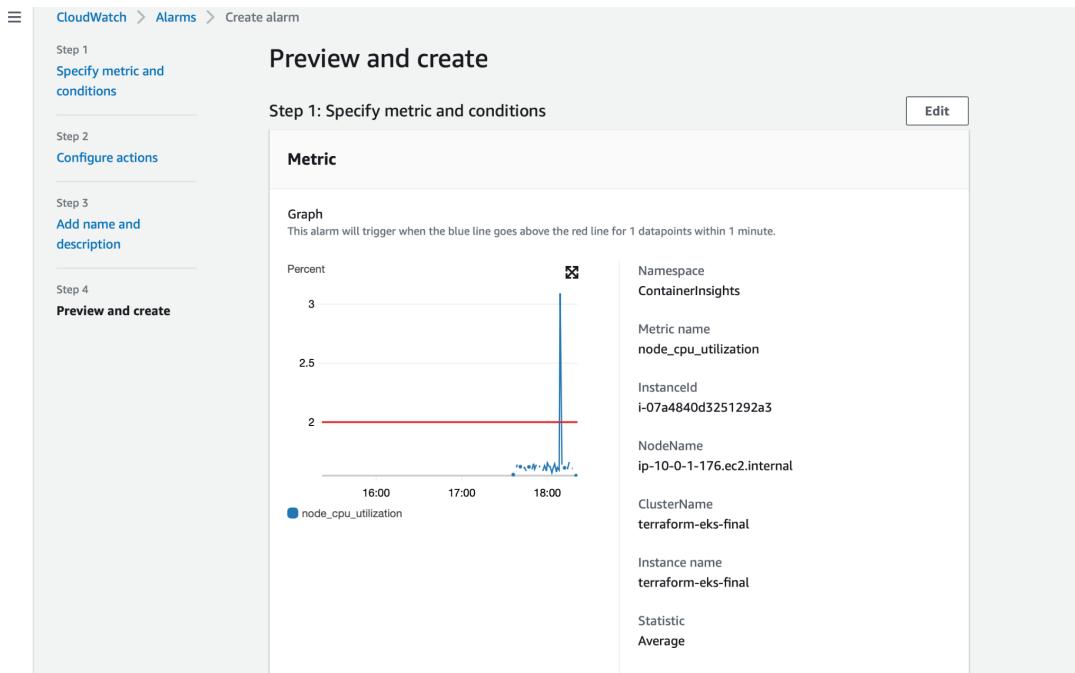
Step 3: Confirm subscription



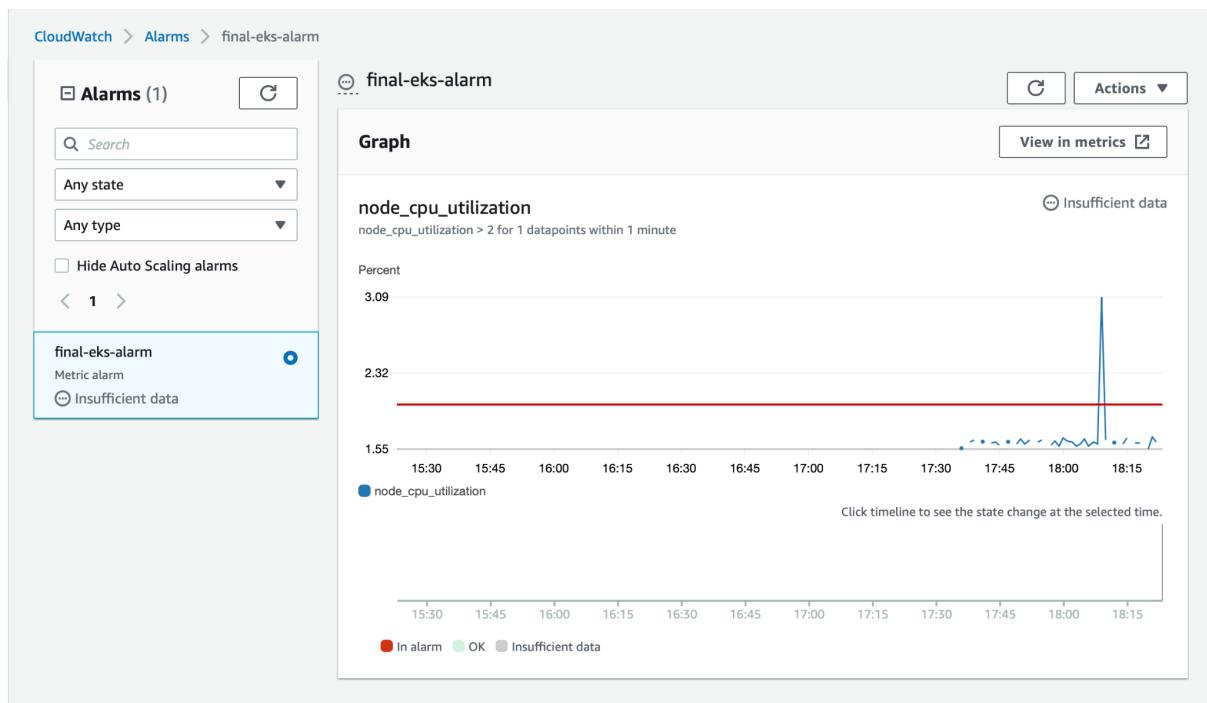
Step 4: Check subscription

```
deepgamit@Deeps-MacBook-Pro EKS % aws sns list-subscriptions-by-topic --topic-arn arn:aws:sns:us-east-1:658346791705:final-
eks-sns
{
    "Subscriptions": [
        {
            "SubscriptionArn": "arn:aws:sns:us-east-1:658346791705:final-eks-sns:7a5fb202-04d8-4749-a9c5-68dc95b2e86b",
            "Owner": "658346791705",
            "Protocol": "email",
            "Endpoint": "deepgamit@gmail.com",
            "TopicArn": "arn:aws:sns:us-east-1:658346791705:final-eks-sns"
        }
    ]
}
deepgamit@Deeps-MacBook-Pro EKS % █
```

Step 5: Setup CloudWatch Alarm to get notified when CPU Utilisation for cluster goes beyond a specific threshold.



Step 6: View CloudWatch Alarm (*final-eks-alarm*) on AWS Console

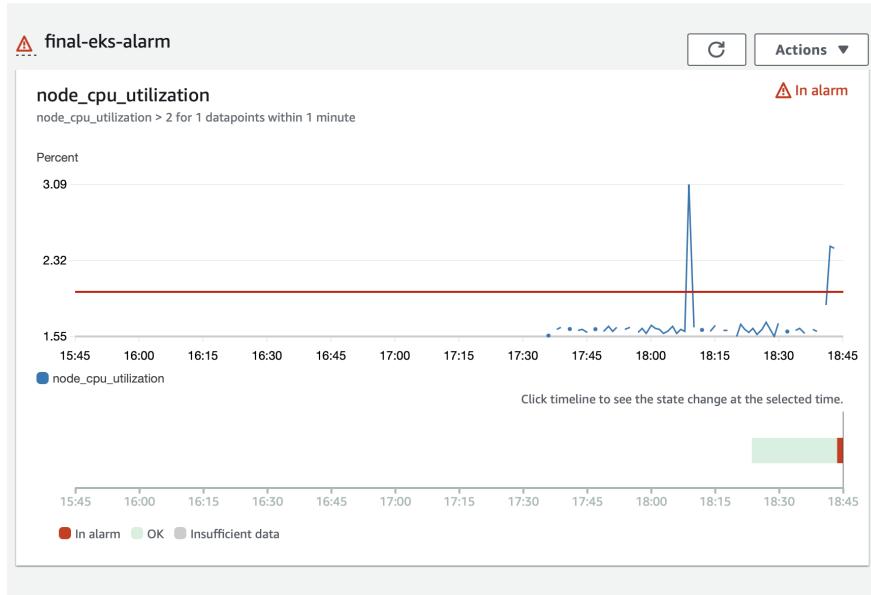


Testing SNS & CloudWatch Alarm

Step 1: Increased load on the app using AB tool

```
deepgmit@Deeps-MacBook-Pro EKS % ab -c 50 -n 5000 http://a3530826f509a499ea27c6fce154b5b-182992700.us-east-1.elb.amazonaws.com/
This is ApacheBench, Version 2.3 <$Revision: 1879490 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
Benchmarking a3530826f509a499ea27c6fce154b5b-182992700.us-east-1.elb.amazonaws.com (be patient)
[...]
```

Step 2: Cloud Watch Alarm Console



Step 3: Check your mailbox for Email

The screenshot shows an email inbox with one new message. The subject of the email is "ALARM: "final-eks-alarm" in US East (N. Virginia)". The email is from "AWS Notifications <no-reply@sns.amazonaws.com> to me". The timestamp is "1:43 PM (1 minute ago)". The email body contains the following text:

You are receiving this email because your Amazon CloudWatch Alarm "final-eks-alarm" in the US East (N. Virginia) region has entered the ALARM state, because "Threshold Crossed: 1 out of the last 1 datapoints [2.4656709987913477 (16/12/21 18:42:00)] was greater than the threshold (2.0) (minimum 1 datapoint for OK -> ALARM transition)." at "Thursday 16 December, 2021 18:43:40 UTC".

View this alarm in the AWS Management Console:
<https://us-east-1.console.aws.amazon.com/cloudwatch/deeplink.js?region=us-east-1#alarmsV2:alarm/final-eks-alarm>

Alarm Details:

- Name: final-eks-alarm
- Description:
- State Change: OK -> ALARM

Reason for State Change: Threshold Crossed: 1 out of the last 1 datapoints [2.4656709987913477 (16/12/21 18:42:00)] was greater than the threshold (2.0) (minimum 1 datapoint for OK -> ALARM transition).

- Timestamp: Thursday 16 December, 2021 18:43:40 UTC
- AWS Account: 658346791705
- Alarm Arn: arn:aws:cloudwatch:us-east-1:658346791705:alarm:final-eks-alarm

Threshold:

- The alarm is in the ALARM state when the metric is GreaterThanThreshold 2.0 for 60 seconds.

Monitored Metric:

- MetricNamespace: ContainerInsights
- MetricName: node_cpu_utilization
- Dimensions: [NodeName = ip-10-0-1-176.ec2.internal] [InstanceId = i-07a4840d3251292a3] [ClusterName = terraform-eks-final]
- Period: 60 seconds
- Statistic: Average
- Unit: not specified
- TreatMissingData: missing

CloudWatch Container Insights

